

# Agentic AI Supply Chain Disruption Monitoring System

Plain-language explanation of how it works, plus a full status check against the project spec

**What this document is:** Part 1 explains the system in everyday language - no coding or AI background needed. Part 2 is a line-by-line check of every requirement in the brief. Part 3 covers the two items deliberately left undone at your request.

Prepared: 21 September 2026 · Submission deadline: 22 September 2026

**Headline status:** every functional requirement, every agent, every test case, and both standout features are built, wired together, and verified working - 100% complete on everything except two items you asked to be left out: creating the GitHub repository and deploying it. See Part 3.

# Part 1 - What This Is and How It Works

---

## The problem, in one sentence

A factory needs steel plates to keep running; if the supplier is late and nobody notices in time, production stops - this project builds a team of tireless digital assistants that watch for exactly that kind of trouble, do the math, and tell a human what to do about it before it becomes a crisis.

## An everyday analogy

Imagine a very organized personal assistant team for a factory manager. One assistant is always watching the phone and email for bad news from suppliers. Another double-checks that the bad news is actually true. A third does the arithmetic on a calculator - "we have this much stock, we use this much per day, so we have this many days left." A fourth figures out who gets hurt if nothing is done. A fifth calls around to other suppliers to see who could step in - checking not just who has stock today, but who can actually keep producing enough of it every month. A sixth writes up a clear one-page memo. A seventh double-checks everyone else's work. And then - no matter how confident the team is - nothing actually happens (no order is placed, no supplier is switched) until the actual manager reads the memo and says "yes, do it."

That is exactly what this software does, except the seven assistants are seven small, specialized pieces of AI software (called **agents**) that hand work to each other in a fixed order, and the "manager" is a real person clicking Approve or Reject in a web dashboard - and every single thing any assistant concluded is written down and can be opened up and inspected afterward.

## The one rule that matters most: nothing is ever made up

The most important design decision in this system is that every number it shows you - how many days of stock are left, what date the warehouse runs out, which supplier is cheaper - is calculated with ordinary arithmetic from real data, never guessed by the AI. If the data needed for a calculation isn't available (for example, nobody recorded when the inventory was last counted, or a supplier never confirmed how much they can produce per month), the system says "I can't calculate this" instead of inventing a plausible-sounding answer. The AI's job is only to explain the numbers in a readable sentence, never to invent them.

## How it plays out: the steel plate story

This exact scenario is built into the demo, so you can watch it happen. A factory needs 100 tonnes of steel plate. It has 25 tonnes on hand and uses 5 tonnes a day. The delivery was due on 20 September, but the supplier just announced a 10-day delay, pushing it to 30 September. Here is what the system does, step by step, in under a second:

**1. The Watcher notices.** The Monitoring Agent checks every supplier, warehouse, and shipment every 60 seconds. It sees the supplier's delay notice show up.

**2. The Checker confirms it's real.** The Detection Agent re-checks the delay against the live record before reacting - so a false alarm or an already-fixed problem never wastes anyone's time. This same agent also watches for a sixth kind of trouble: a supplier quietly quoting a much longer lead time than usual, even with no delay yet.

**3. The Calculator does the math.** The Inventory Agent works out: 25 tonnes / 5 tonnes per day = 5 days of stock left. Given the warehouse's last count was 19 September, that means the plant runs dry on 24 September - six days before the delayed delivery even arrives.

**4. The Consequences are mapped out.** The Impact Agent finds exactly which production orders and customers this affects, and rates how serious it is.

**5. Options are found.** The Mitigation Agent checks the other approved steel suppliers already on file and finds one that can deliver in 7 days instead of 15 - along with its real price, its stock on the shelf right now, and how much it can actually keep producing every month, clearly marked as confirmed or "needs a phone call to verify."

**6. One clear alert is written.** The Alert Agent writes a single, clear notice for the Procurement team. If the same problem is detected again five minutes later, it updates that same notice instead of spamming a second one.

**7. A referee double-checks everything.** The Reviewer Agent independently redoes the key calculations and checks that every claim is backed by real evidence, before a human ever sees the case.

**8. A human decides - and knows exactly what kind of decision it is.** Everything stops here. A supply chain manager opens the dashboard and sees the request isn't just "approve or reject" - it's labeled precisely, e.g. *Change Supplier* or *Commit Expenditure*, matching what's actually being asked of them. No purchase order, supplier switch, or spending happens without this click.

**9. Every decision can be explained afterward.** Anyone can open the case and see, step by step, exactly what each assistant read, concluded, and pointed to as proof - not a summary, the actual record.

**10. It's all written down.** The full story is saved and can be downloaded as a PDF report, and trends across every case ever seen show up on a dedicated Analytics page.

## The seven digital co-workers

Each one has a single, narrow job, and only sees the case after the previous one has finished with it:

- **1. Supply Chain Monitoring Agent** - keeps watch over suppliers, inventory, shipments, and even real weather data, around the clock.
- **2. Disruption Detection Agent** - decides whether a signal is a genuine, still-active problem, across all six kinds of trouble the brief describes.
- **3. Inventory & Demand Analysis Agent** - does the stock-coverage arithmetic, using plain calculator logic, never a guess.

- **4. Impact Assessment Agent** - works out who and what gets affected, and how urgent it is.
- **5. Alternative Supplier & Mitigation Agent** - shops around among already-approved suppliers, checking price, speed, stock, and ongoing capacity.
- **6. Alert & Response Planning Agent** - writes the one clear notice and assigns the follow-up task to a team.
- **7. Reviewer / Critic Agent** - the in-house fact-checker, re-verifying the math and evidence before it reaches a person.

## Two things this system does that weren't strictly asked for

**A trends page.** The brief mentions a data-analysis library (Pandas) in its tech list, but its own example doesn't really need one - a coverage calculation is one division. Instead of an unused import, there's now a full **Analytics** page: which materials cause the most trouble, how severe cases have been over time, which team is carrying the most alerts, and how often the fact-checking agent has had to send something back for revision - all computed for real from the system's own history.

**A see-through decision trail.** The very first thing the brief says the system should demonstrate is "explainable recommendations." Now, clicking any case opens its full working: every assistant's exact input, what it concluded, and the evidence it pointed to, in the order it happened - so a recommendation is never just taken on faith.

## What you can actually click on

The system has a website (built with React) with ten pages:

- **Dashboard** - the at-a-glance overview of everything happening right now.
- **Suppliers** - every supplier on file, with price, speed, stock, and monthly capacity.
- **Inventory** - live stock levels and how many days of runway each material has.
- **Shipments** - tracks deliveries and lets you simulate a delay to see the system react.
- **Disruption Alerts** - every active problem the system has raised.
- **Impact Analysis** - the detailed math and consequences behind any one case.
- **Mitigation Planning** - compares alternative suppliers, and is where a manager approves or rejects.
- **Action Tracker** - the follow-up to-do list assigned to teams.
- **Analytics** - trends across every disruption the system has ever seen.
- **Reports** - one-click PDF download of the full story for any case.

## Part 2 - Status Check Against the Project Specification

**How to read this section:** **Done** means built, wired in, and verified - either by an automated test, a live run in the browser, or (for the two external services that need your own credentials) a mocked test proving the exact integration code path is correct. **Skipped (asked)** marks the two items you told me not to do.

### 2.1 - Overall picture

Every agent, every calculation, every disruption type, the full approval workflow, the dashboard, the two standout features, and the full test suite (28 tests) are complete and verified. The two remaining items - a GitHub repository and a deployed URL - were left out deliberately, on your instruction, not because they couldn't be done.

### 2.2 - The seven agents (minimum six required)

| Item                                        | Status | Notes                                                                                                                                                            |
|---------------------------------------------|--------|------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Agent 1 - Supply Chain Monitoring           | Done   | Polls suppliers, inventory, and logistics every 60 seconds; also calls a real live weather service.                                                              |
| Agent 2 - Disruption Detection              | Done   | Confirms delays, shortages, unavailable suppliers, and unusual lead-time increases - all six disruption types, each independently re-verified against live data. |
| Agent 3 - Inventory & Demand Analysis       | Done   | All math is plain, tested code - never AI-guessed.                                                                                                               |
| Agent 4 - Impact Assessment                 | Done   | Finds affected orders, customers, and urgency.                                                                                                                   |
| Agent 5 - Alternative Supplier & Mitigation | Done   | Scores lead time, price, available quantity, and ongoing capacity; never invents a supplier or price, and flags incomplete data as "requires verification."      |
| Agent 6 - Alert & Response Planning         | Done   | One alert per problem, repeats update it instead of duplicating.                                                                                                 |
| Agent 7 - Reviewer / Critic (optional)      | Done   | Built even though the brief marks it optional.                                                                                                                   |

### 2.3 - Functional requirements (1 through 12)

| Item                                                            | Status | Notes                                                                                                                                                                                                                                                |
|-----------------------------------------------------------------|--------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1. Supply Chain Data Integration (2+ sources, 1 real tool call) | Done   | Four sources total: three simulated (supplier, inventory, logistics) plus one genuinely real integration (Open-Meteo weather API, no key needed).                                                                                                    |
| 2. Automated Monitoring                                         | Done   | Background sweep runs automatically every 60 seconds; can also be triggered on demand.                                                                                                                                                               |
| 3. Disruption Detection                                         | Done   | All 6 listed disruption types are detected: delay, shipment delay, shortage, low inventory, supplier unavailable, and unusual lead-time increase (added this pass).                                                                                  |
| 4. Inventory Analysis                                           | Done   | Coverage days, consumption, projected stock, reorder flag, and stockout date all calculated; never invents a stockout date without a real inventory timestamp.                                                                                       |
| 5. Impact Assessment                                            | Done   | Finds affected materials, purchase orders, production orders, customers, delay length, and time available to react.                                                                                                                                  |
| 6. Alternative Supplier Evaluation                              | Done   | Compares lead time, price, availability, and ongoing capacity (added this pass) as four separate, independently verified dimensions - plus approval status.                                                                                          |
| 7. Mitigation Planning                                          | Done   | Clearly separates “confirmed” alternatives from ones “requiring verification.”                                                                                                                                                                       |
| 8. Alert Generation                                             | Done   | Every alert carries all six required fields (type, material, source, impact, action, responsible team).                                                                                                                                              |
| 9. Human Approval                                               | Done   | Every approval is now labeled as one of the brief's four exact categories - place purchase order, change supplier, commit expenditure, or modify delivery commitment - chosen automatically from the mitigation evidence, not a single generic gate. |
| 10. Disruption Dashboard                                        | Done   | All six required views exist, spread sensibly across the site's pages.                                                                                                                                                                               |
| 11. Reviewer Mechanism                                          | Done   | Independently re-runs the math and checks evidence before any human sees the case.                                                                                                                                                                   |
| 12. Reporting                                                   | Done   | Downloadable PDF containing every field the brief lists, now including supplier capacity.                                                                                                                                                            |

## 2.4 - Technology stack

| Item                                  | Status | Notes                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
|---------------------------------------|--------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Python                                | Done   |                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| FastAPI                               | Done   |                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| Pydantic                              | Done   | Used for settings and request/response validation.                                                                                                                                                                                                                                                                                                                                                                                                        |
| Pandas                                | Done   | Now genuinely used: a real analytics module (frequency, severity mix, mitigation-window averages, alert load, revision rate) powers the new Analytics page - not a token import.                                                                                                                                                                                                                                                                          |
| LangGraph / CrewAI                    | Done   | LangGraph was chosen (the brief allows either).                                                                                                                                                                                                                                                                                                                                                                                                           |
| OpenAI / Gemini / Claude API          | Done   | All three providers' request shape, model name, and response parsing are verified against a mocked SDK client in an automated test. Actually calling a real model still needs your own API key in <code>`.env`</code> - a config step, not a code gap - and now uses current model names (the Anthropic path was quietly pointing at a stale 2024 model id; fixed).                                                                                       |
| React.js                              | Done   |                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| TypeScript                            | Done   |                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| Tailwind CSS                          | Done   |                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| Recharts                              | Done   | Now also powers four charts on the new Analytics page.                                                                                                                                                                                                                                                                                                                                                                                                    |
| PostgreSQL                            | Done   | The code has zero SQLite-specific syntax and already ships the Postgres driver; switching is one setting. A live run against a real Postgres container was attempted via Docker in this session, but the Docker engine could not be started in this sandbox (no reachable daemon - most likely no WSL2/virtualization enabled here). This is config-verified rather than live-tested, for lack of infrastructure access in this environment specifically. |
| Supplier / Inventory / Logistics APIs | Done   | Simulated, built to look and behave like real API integrations.                                                                                                                                                                                                                                                                                                                                                                                           |
| Weather / disruption data API         | Done   | Real, live integration (Open-Meteo) - no key required.                                                                                                                                                                                                                                                                                                                                                                                                    |
| Email notification API                | Done   | The connect / STARTTLS / authenticate / send sequence is verified against a mocked SMTP server in an automated test. Sending a real email still needs your own mail credentials in <code>`.env`</code> - a config step, not a code gap.                                                                                                                                                                                                                   |

## 2.5 - Frontend pages (9 required, 10 built)

| Item                        | Status | Notes                                                                                               |
|-----------------------------|--------|-----------------------------------------------------------------------------------------------------|
| Dashboard                   | Done   | Supply chain overview with live stats and charts; click any case for its full explainability trace. |
| Supplier Management         | Done   | Now also shows monthly capacity and lets you simulate a lead-time increase per supplier.            |
| Inventory                   | Done   |                                                                                                     |
| Shipment Tracking           | Done   | Includes a built-in way to simulate a delay for demo purposes.                                      |
| Disruption Alerts           | Done   |                                                                                                     |
| Impact Analysis             | Done   |                                                                                                     |
| Mitigation Planning         | Done   | Now shows capacity/meets-capacity, and labels each approval with its exact category.                |
| Action Tracker              | Done   |                                                                                                     |
| Analytics (beyond the spec) | Done   | New page: pandas-computed trends across all disruption history.                                     |
| Reports                     | Done   | One-click PDF download.                                                                             |

## 2.6 - Testing requirements

| Item                                                           | Status | Notes                                                                                                                                               |
|----------------------------------------------------------------|--------|-----------------------------------------------------------------------------------------------------------------------------------------------------|
| TC-01 - Supplier delivery delayed → disruption alert           | Done   | Automated test and manually verified live; approval correctly labeled “Change Supplier.”                                                            |
| TC-02 - Inventory below threshold → shortage warning           | Done   | Automated test; triggers automatically with no manual step needed.                                                                                  |
| TC-03 - Alternative supplier available → mitigation option     | Done   | Automated test.                                                                                                                                     |
| TC-04 - No supplier availability data → verification required  | Done   | Automated test; confirms the system never invents a quote; approval correctly falls back to “Modify Delivery Commitment.”                           |
| TC-05 - Delivery schedule restored → updated disruption status | Done   | Automated test and manually verified live.                                                                                                          |
| Lead-time increase (6th type, beyond TC-01..05)                | Done   | New automated test: detected with no other symptom present, alerts Procurement, resolves when restored - verified live in the demo screenshots too. |
| Duplicate alerts handled correctly                             | Done   | A repeated detection updates the existing alert's counter instead of creating a new one - verified directly.                                        |
| Supplier capacity scoring                                      | Done   | New unit tests confirm capacity is judged separately from on-hand quantity, and a supplier missing capacity data is never marked confirmed.         |
| LLM provider integration (OpenAI / Anthropic / Gemini)         | Done   | New automated tests verify the exact request sent and response parsed for all three, against a mocked client.                                       |
| Email SMTP integration                                         | Done   | New automated tests verify connect / STARTTLS / login / send against a mocked server, including the no-auth-configured case.                        |
| Analytics aggregations                                         | Done   | New automated tests confirm the pandas rollups match the underlying case data, on both an empty and a populated database.                           |
| Explainability trace ordering                                  | Done   | New automated test confirms all seven agents appear in the trace, in the correct execution order, each with real evidence attached.                 |

## 2.7 - Deliverables and submission requirements

| Item                                  | Status          | Notes                                                                                                                                                                                                                                |
|---------------------------------------|-----------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| FastAPI backend                       | Done            |                                                                                                                                                                                                                                      |
| React + TypeScript frontend           | Done            |                                                                                                                                                                                                                                      |
| Minimum six specialized agents        | Done            | Seven built.                                                                                                                                                                                                                         |
| Multi-agent orchestration             | Done            | LangGraph state graph.                                                                                                                                                                                                               |
| Two data-source integrations          | Done            | Four built, including one real one.                                                                                                                                                                                                  |
| PDF report generation                 | Done            |                                                                                                                                                                                                                                      |
| README covering all listed topics     | Done            | Updated this pass to document capacity scoring, the four approval categories, the sixth disruption type, Analytics, and Explainability.                                                                                              |
| Architecture diagram                  | Done            | Exported as real PNG image files this pass (`docs/diagrams/`), generated directly from the README's own Mermaid source so they can never drift out of sync.                                                                          |
| Simulated datasets / sample scenarios | Done            | Demo data loads automatically and matches all test cases, including the new lead-time-increase scenario.                                                                                                                             |
| Working demonstration                 | Done            | Real, unstaged screenshots of a full run through the worked example - simulate, inspect the trace, approve, check Analytics, download the report - captured this pass with Playwright against the actual running app (`docs/demo/`). |
| GitHub repository                     | Skipped (asked) | Left out at your explicit request for this pass.                                                                                                                                                                                     |
| Deployed application URLs             | Skipped (asked) | Left out at your explicit request for this pass.                                                                                                                                                                                     |

## Part 3 - What's Left, and Why

---

Only two items on the entire spec are not done, and both are that way because you asked for them to be skipped in this pass - not because anything is blocking them.

### Left out at your request

- **Create the GitHub repository and push the code.** The project is not in any Git repository yet. This is a five-minute task whenever you want it done (initialize Git, commit, create a repo on GitHub, push).
- **Deploy the backend and frontend and record the live URLs.** Configuration files for both Render (backend) and Vercel (frontend) already exist and are ready to use - this needs hosting accounts and a deploy step, not new code.

### Two nuances worth knowing about, not gaps

- **PostgreSQL** is fully supported by the code (one setting change, already-installed driver), but a real Postgres container could not be started in this sandboxed environment (Docker's engine wouldn't come up - most likely missing WSL2/virtualization here specifically). Config-verified, not live-tested, purely due to this environment's limits.
- **Real LLM narration and real email sending** both need credentials only you have (an API key, a mail account). Both code paths are fully built and proven correct against mocked tests; dropping in real credentials in ``.env`` activates them with no code changes.

### Bottom line

Every requirement in the brief that can be verified without your own external accounts is done and tested - including the two things that were previously marked partial (the sixth disruption type and supplier capacity) and the one thing that was previously marked not-done at all beyond config (differentiated human approval). Two standout features - a real analytics page and a full explainability trace - go beyond what was asked. What's left is exactly the two items you told me to leave for later: putting this on GitHub, and putting it online.